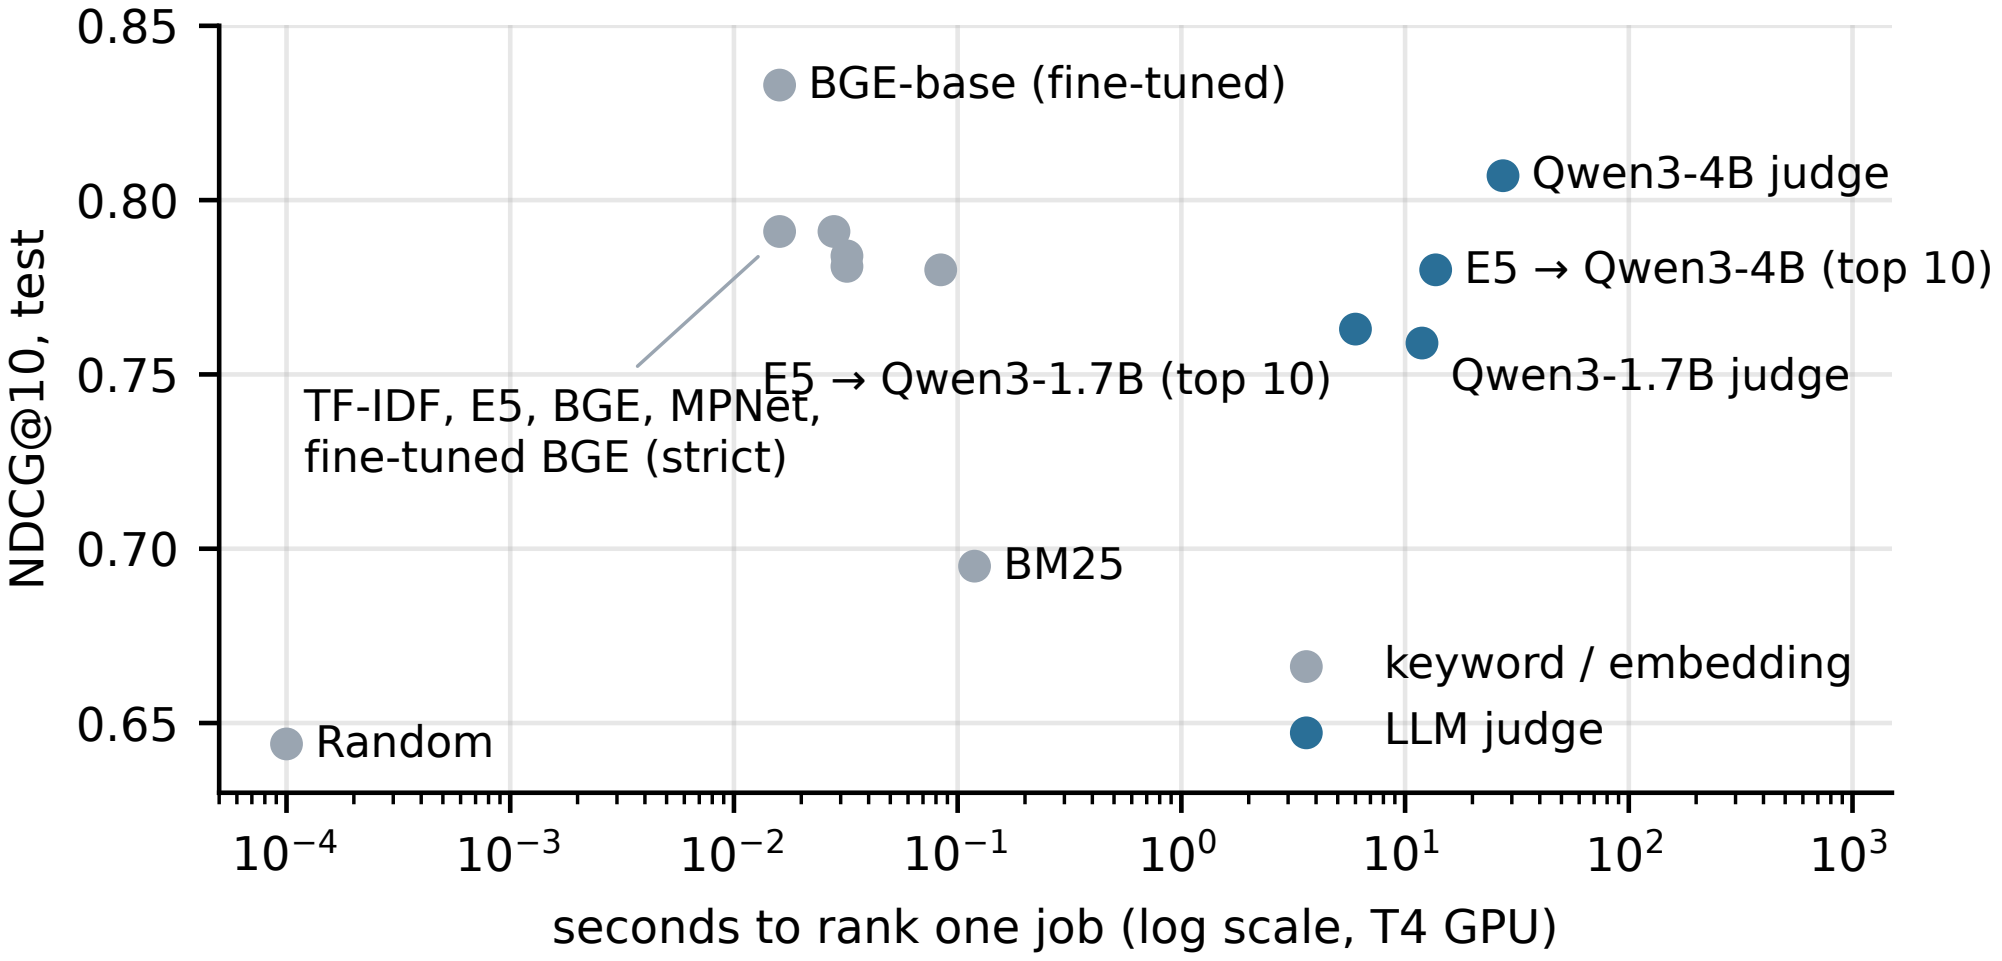
0.85
0.80
0.75
0.70
0.65
NDCG@10, test
BGE-base (fine-tuned)
Qwen3-4B judge
E5 → Qwen3-4B (top 10)
E5 → Qwen3-1.7B (top 10)
Qwen3-1.7B judge
TF-IDF, E5, BGE, MPNet,
fine-tuned BGE (strict)
BM25
keyword / embedding
LLM judge
Random
10^{-4}
10^{-3}
10^{-2}
10^{-1}
10^{0}
10^{1}
10^{2}
10^{3}
seconds to rank one job (log scale, T4 GPU)